
Twitter or WhatsApp?

Public and Private Agent Networks Are Different at the Edge

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Agent networks are being built to look like Twitter or like WhatsApp, and the two
2 are compared as though “public versus private” were a variable. It is not. An edge
3 is a triple—who may form it, what it carries once formed, what is merely asserted
4 about it—and the two regimes are not two values of one dial but two *bundles*
5 that differ on all three at once: lateral breadth, sandboxed exchange and platform-
6 computed scores against longitudinal depth, openable state and relationship-carried
7 history. Comparing them measures a bundle swap, and which number it returns
8 is set by a pairing nobody documents: the two diagonals of the same factorial
9 give 0.181 and 0.314 on the same episodes. We give an instrument that sets the
10 three independently—one model, one runtime, one coordination kernel, conditions
11 expressed as capability grants—and a metric set that prices each axis in accuracy,
12 tokens and disclosure rather than in accuracy alone. Across 582 episodes of
13 hidden-profile coordination scored against forced-value answer keys, the factorial
14 recovers what neither diagonal isolates—0.248, 0.066 and 0 on requirement F_1 :
15 a length-matched placebo moves the score as much as a trust assertion does, and
16 the formation effect is reach, vanishing when the roster already holds what a task
17 needs. Two results were not designed. A bounded roster is shown a quarter as much
18 misinformation as an open directory and writes more of it into the deliverable,
19 fabricating values 58 times where the open arms fabricated once—filtering out bad
20 claims also removes the second source that catches them. And a single brokered
21 hop recovers +0.140 of the gap while every assertion of trust remains flat, which
22 separates an introduction, a formation intervention, from an endorsement, an
23 attribute. The weights are measured on one task family; whether they are constants
24 of the framework or of the task is the experiment the framework now specifies.

1 Introduction

26 Two pictures of an agent network are in circulation. In the first, agents discover each other in the
27 open: they publish what they can do, broadcast what they need, and are matched by relevance rather
28 than acquaintance. In the second, agents act for people who already work together: membership is
29 bounded, interaction repeats, context accumulates. The first is called public and the second private,
30 and the difference is usually described as one of visibility.

31 Visibility is the wrong frame, and so is treating the two as endpoints of one dial. On Twitter an
32 account addresses any other, relevance is a claim in a profile, and a bad actor changes handles. On
33 WhatsApp a conversation exists because a number was already held or someone made an introduction,
34 the tie carries files and history, and misbehaviour costs the tie and embarrasses whoever vouched.

35 The two differ in what it costs to open an edge, in what the edge carries once open, and in what is
36 asserted about the counterpart—three things at once, not one.

37 Take the edge as the primitive and the three separate. **Edge formation** is which edges a requester
38 may create at all. **Edge semantics** is what an existing edge carries: a delegated task returning a result,
39 or direct access to state. **Edge attributes** are the beliefs attached to an edge—a reputation score, a
40 stated shared objective, a record that the parties have worked together before—which change nothing
41 a runtime checks. Public and private are two *bundles* in this space, not two definitions and not two
42 settings of one dial, so a study that compares them swaps all three at once. Worse, the two ways of
43 pairing the axes give different answers on the same data—0.181 and 0.314—so the reported size of
44 the “public advantage” is set by a bundling choice, not by a measurement. **Public versus private is**
45 **not the experimental variable**; $do(F)$, $do(S)$ and $do(A)$ are.

46 **This paper unbundles them and prices the components.** We place the two regimes as corners of
47 (F, S, A) rather than as definitions (§2). We give an instrument that sets the three axes independently,
48 and a metric set that charges each one in accuracy, tokens and disclosure (§3). And we state the design
49 question the decomposition makes answerable: not which regime wins, but which configuration a
50 setting should hold, and whether the axis weights are constants of the framework or of the task (§4).
51 The coordination kernel is an instrument, not a contribution.

52 **Contributions.** (i) **Public and private recast as two bundles in an edge space** $e = (F, S, A)$
53 rather than as two definitions, which makes “public versus private” a bundle swap rather than a
54 variable (§2). (ii) **An instrument that sets the three axes independently**—one model, one runtime,
55 one kernel, conditions as capability grants—with a metric set that prices an axis in accuracy, tokens
56 and disclosure rather than accuracy alone (§3). (iii) **A configuration problem in place of a binary**
57 **comparison**, whose first result is that the space is smaller than it looks and whose sharpest untested
58 prediction is that an edge’s position moves (§4). (iv) **A first decomposition on one task family**, with
59 a length-matched placebo on the attribute axis and two undesigned results (§5).

60 We do not study how a stranger becomes a teammate, and we make no safety claim about either
61 regime. We state in advance that the *existence* of a task-dependent crossing between regimes follows
62 from how we allocate facts and is not a discovery; what is measured is where it falls and what it
63 costs.

64 2 Public and private are two corners, not two definitions

65 Take the network as the unit of analysis: agents are nodes, a request from one to another is an edge.
66 An edge is not one thing. It is a triple

$$e = (F, S, A)$$

67 — **Formation**, who may open it; **Semantics**, what it carries once open; **Attributes**, what is asserted
68 about it. The three are independently setttable, and the first two are checked by a coordination layer
69 on every call while the third only ever reaches a model as text.

70 **Public and private are not points on one dial, and they are not definitions we**
71 **supply. They are two characteristic bundles in this space** — two corners that
72 co-occur in deployed systems because each component makes the others affordable.
73 Every real network sits somewhere between them, and most sit at no corner at all.

74 Twitter and WhatsApp are the two corners. On Twitter an account addresses any other without
75 assent, relevance is a claim in a profile, exchanges reveal nothing but their content, and trust arrives
76 as a platform-computed badge; a bad actor changes handles. On WhatsApp a conversation exists
77 because a number was already held or someone made an introduction, the tie carries files and history,
78 what is known about the counterpart is who introduced them and how the last exchange went, and
79 misbehaviour costs the tie and embarrasses whoever vouched.

80 The consequence is the paper’s premise:

81 **“Public versus private” is not an experimental variable.** It is a comparison
82 between two bundles that differ on three axes at once. A study that runs it measures

Table 1: The three axes, and what each regime supplies on it. Formation and semantics are enforced by a coordination layer on every call; attributes only ever reach a model as text. A study that compares a public network to a private one moves all three.

	Public	Private
Edge formation	Unilateral. Anyone in the directory is addressable; matching is by declared relevance. Lateral topology : wide and shallow, many weak ties.	Bilateral or brokered. A counterpart is reachable because of a prior tie or an introduction. Longitudinal topology : narrow and deep, few strong ties.
Edge semantics	Sandboxed by default. A counterpart accepts a task and returns a result; its store stays opaque. Interaction is one-shot request/response.	Openable. Shared workspaces, persistent namespaces, a store that may be listed and read—affordable precisely because the tie exists.
Edge attributes	Platform-computed for strangers: reputation scores, ratings, usage counts, verification badges.	Carried by the relationship: tenure, who introduced whom, how the last collaboration went, who is answerable.

83 $(F_{\text{pub}}, S_{\text{pub}}, A_{\text{pub}}) \rightarrow (F_{\text{pri}}, S_{\text{pri}}, A_{\text{pri}})$ and reports one number, which cannot be
84 attributed to any component. The variables are $do(F)$, $do(S)$ and $do(A)$, one at a
85 time.

86 Table 1 is what each corner supplies on each axis — the bundles unpacked.

87 The rest of this section takes them one at a time.

88 **The two topologies are the substance of the formation axis.** A public directory buys *breadth*:
89 a rare requirement finds someone and—less obviously—a claim can be put to a second source. A
90 private roster buys *depth*: fewer counterparts, each already paid for, each with a history. These are
91 not two settings of one dial; they are the returns on two different edge prices. Which links form, and
92 at what price, is strategic network formation [Jackson and Wolinsky, 1996], and the agent setting
93 differs in that the price is conventionally zero. On *semantics*, exposure follows the tie—one does not
94 open a working directory to a stranger—but the axis is separately settable, and a capability kernel
95 expresses both settings as grants over different namespaces, so a condition is a set of capabilities
96 rather than a different codebase [Google]. On *attributes*, both regimes carry a characteristic form of
97 trust-conveying metadata that no runtime checks: a public network computes a score about a stranger,
98 a private one carries a history about an acquaintance. Both rest on the same assumption—that an
99 agent builds a representation of its counterpart from what it is told and acts differently because of it.

100 2.1 Why these three and not nine

101 The literature on ties, teams and open innovation supplies at least nine relational dimensions—
102 membership openness, identity continuity, shared-state depth, repeated interaction, relationship
103 origin, prior trust, shared goals, accountability, exit cost. They are not nine degrees of freedom: in a
104 deployed system an open directory brings weak identity, which brings diffuse accountability, which
105 brings low exit cost. The three-way grouping breaks that coupling the one way a controlled study can.
106 Formation and semantics are properties of the interaction layer—independently settable, checked
107 on every call, not open to argument by a model. Attributes are properties of the people behind the
108 agents; a harness cannot instantiate them, only *assert* them. That asymmetry is the experiment rather
109 than a limitation to apologise for: deployed networks communicate trust, alignment, attribution and
110 tenure as text on cards, so asking whether a model acts on that text asks whether a large part of
111 agent-network design does anything at all.

112 2.2 Two corrections the framework has to carry

113 Both fall out of taking the edge, rather than visibility, as the primitive. First, **a public network’s**
114 **noise is partly a feature**: a requester who can address many counterparts can put the same question
115 to a second one, and contradiction is *detectable*, so a roster that filters misinformation out before it
116 arrives also removes the second source needed to catch what does arrive. Whether the filter or the
117 redundancy dominates is empirical, and the framework must not settle it by calling breadth a cost.

118 Second, **an endorsement is not an introduction**, and they sit on different axes. “This counterpart is
 119 **reliable**” is an attribute: text on an edge that already exists. An introduction that makes a third agent
 120 **addressable** is formation: it changes the graph. A framework that files both under “trust” predicts
 121 that a private network’s advantage comes from what it knows about its members; the decomposition
 122 predicts something narrower—that only the half which changes reachability can do work—and §5
 123 reports which way it went.

124 3 Making each axis measurable

125 The framework is only useful if each axis can be set independently and each regime’s claimed
 126 advantage can be priced. Model, runtime, task, answer key and coordination kernel are held fixed;
 127 a condition is a set of capability grants rather than a different codebase, so **formation** (which cards
 128 a requester may enumerate and address), **semantics** (whether a counterpart returns an answer or
 129 exposes a store) and **attributes** (what the cards say about themselves and each other) move one at a
 130 time. Formation and semantics are refused by the kernel when violated; attributes reach the model
 131 only as tokens, which is what makes the third axis a test rather than a tautology. Appendix B gives
 132 the instrument.

133 The task class has to force pooling, or reach cannot be worth anything. Each episode is a hidden-
 134 profile deliverable: the facts a correct artifact needs are distributed across counterparts so that no
 135 single agent, including the requester, holds enough [Stasser and Titus, 1985]. E, the share of required
 136 facts placed *outside* the roster, is the dose on the formation axis—at $E = 0$ a private network needs
 137 nothing it cannot already reach, so the axis must show no effect there. Distractors plant contradicting
 138 values on counterparts other than the true holder, so that a second source has something to catch.
 139 Every fact is a number, a date, a name or an enum, and the deliverable is scored by exact match after
 140 a normalization rule fixed before the keys were written: no judge sits on any primary metric.

141 **An axis is priced, not just scored.** Around requirement F_1 sit useful sources (counterparts
 142 contacted that held a needed fact—the mechanism variable for formation), **pollution** absorbed
 143 and invented (a distractor’s value appearing in the artifact, split by whether the agent was ever
 144 shown it, the second being fabrication), and **tokens and contacts**—the price of the arm and, since
 145 each contact discloses the goal, its exposure. **Semantics** buys accuracy and is charged in tokens;
 146 reporting only the first half would make the private regime look better than the exchange rate supports.

147 **Delivery is reported before any contrast.** A trial that produced no artifact is a missing observation,
 148 not a score of zero, and averaging it in as zero converts an outage into an effect. We found this the
 149 expensive way (§A.7). Every contrast we report is therefore preceded by a per-cell delivery table, is
 150 computed both conditional on delivery and with losses scored as zero, and is treated as uninterpretable
 151 when delivery differs across compared cells by more than 15 points.

152 4 From a binary comparison to a configuration problem

153 Once the bundle is unpacked, the question changes shape. It is no longer which corner wins; it is
 154 which point in (F, S, A) a given setting should hold. Three binary axes give eight configurations,
 155 and no deployed network is a corner of them: a roster with an introduction protocol is public on
 156 formation for whoever a broker can reach and private for everyone else; an enterprise workspace is
 157 private on formation and fully open on semantics.

158 **The first result of the decomposition is that the space is smaller than it looks.** A configuration
 159 search is only a search if more than one dimension moves the outcome. On the task family we ran,
 160 the three weights are 0.248, 0.066 and 0 on requirement F_1 (§5), and the third is not small—it is
 161 indistinguishable from two neutral lines of text. **An axis with zero weight is not a dimension.** Eight
 162 configurations collapse to four. Semantics survives—0.066, interval $[+0.012, +0.119]$ —but it is
 163 charged at 11k tokens an episode, so on this task family the four collapse in practice to a choice on
 164 formation with a priced option attached.

165 That collapse is a finding, not a failure of the framework, and it is the kind of finding only a
 166 decomposition can produce: a study that compared the two corners would have reported one diagonal—
 167 0.181 or 0.314, depending on which pairing it happened to call public and private—and had no way to

168 learn that one third of the bundle does nothing. We report it rather than restoring the eight-way search,
169 because the honest next question is sharper than an optimisation. **Is the space two-dimensional**
170 **everywhere, or only here?**

171 **Scenarios are the test of whether the weights are constants.** The collapse above was measured
172 on a single task class: facts held by different agents, aggregated into one deliverable. That class is
173 maximally favourable to formation—reach is literally what it asks for—and maximally unfavourable
174 to semantics, since there is no shared artifact to co-edit. If the three weights are properties of the
175 framework, they should hold elsewhere. If they are properties of the task, they will move, and the
176 combination problem becomes real.

177 The enterprise row is the sharpest test, because it predicts the reverse of what we measured: when
178 everyone needed is already reachable, formation has no work left to do, and the shared store that
179 cost 11k tokens for a small gain in a pooling task should pay for itself when the deliverable is edited
180 jointly. It is also the row we defer—it needs a co-edited task family rather than a reallocation of
181 ours, and we are pursuing it as dedicated work at a venue that can give the setting the space it needs.
182 A framework that only ever produces “formation dominates” is a restatement of the task; one that
183 predicts where formation stops dominating, and is then held to it, is a framework.

184 4.1 Edges may start public and end private

185 The framework’s sharpest prediction is not about a static configuration. If the axes are separable,
186 an edge’s position in (F, S, A) can *move*, and each axis unlocks the next: open discovery supplies
187 the second source, a sandboxed first exchange makes a stranger affordable, repeated success is the
188 only thing that could make an attribute mean anything—a reputation mark is inert here precisely
189 because nothing generated it—and shared state is worth 11k tokens only once an edge expects to be
190 reused. **We have not tested it.** Every episode we ran is a cold start: 582 of them with the repeat knob
191 present in the harness and never turned. If an open requester’s contact set narrows across repeats,
192 “public versus private” is a question about *time* rather than architecture, and a private network is what
193 a public one becomes; if it does not, the transition must be designed rather than left to emerge.

194 **What an optimum means here.** An optimum is a configuration, not a size. On formation it is a
195 *funnel depth*—own roster, one brokered hop, two hops, open directory—where each layer widens
196 reach and weakens recourse. The marginal return on a layer is the number of *independent* sources it
197 adds, which saturates; the marginal cost is what a claim from that layer is worth, which decays. The
198 crossing is the depth, and per §4.1 it may be a schedule rather than a constant.

199 5 Evidence so far

200 We have run the decomposition on one task family: 582 episodes across sixteen sweeps, one model
201 family, one kernel. The full record—every result, the nulls, and one retraction—is Appendix A. Four
202 findings bear on the framework.




Figure 1: The formation effect is reach, and the dose-response says so. E is the share of a task’s required facts placed outside the roster; bars are 95% percentile bootstrap intervals over episodes. Four things hold at once, and the axis needs all four: the contrast **vanishes at zero dose**, which is the control it had to pass; the arm that cannot be affected is **flat** across all three doses; the mechanism variable tracks the dose (the bounded arm’s useful sources fall $4.36 \rightarrow 3.17 \rightarrow 1.26$); and fabrication appears only as sources dry up ($0 \rightarrow 0.139 \rightarrow 0.333$ per task).

Finding	Evidence
Only formation is an intervention	The two diagonals of the same factorial give 0.181 and 0.314; the factorial recovers 0.248 formation and 0.066 (CI [+0.012, +0.119], +11k tokens) semantics, with ≈ 0 on attributes. 56 contrasts per beat against a chance floor of 2.8; the two that cleared reverse sign between beats of the same episodes; a length-matched placebo of two neutral lines moves the score +0.005.
The formation effect is reach	Dose-response in E (Figure 1): +0.034 (CI [−0.067, +0.137]) at $E = 0$, +0.111 at 0.3, +0.279 at 0.7. The public arm is flat across all three doses; the private arm’s useful sources fall $4.36 \rightarrow 3.17 \rightarrow 1.26$.
A roster sees less bad information and believes more of it	Shown a third as much planted misinformation; writes 0.269 against 0.094 instances per task into the deliverable (+0.175, CI [+0.105, +0.242]). Across 287 and 286 delivered tasks the public arms fabricated once, the private arms 58 times. Matching on sources obtained does not close it (2.4% against 18.1% at four sources).
An introduction does work an endorsement does not	One brokered hop recovers +0.140 (CI [+0.073, +0.204]) of a 0.304 gap, replicating within three edge-cost strata, and cuts fabrication $0.223 \rightarrow 0.057$. Read-all with deep ties only inside the roster is indistinguishable from full openness (+0.019). Both are formation manipulations; every attribute manipulation is flat.

What this does not cover. One model family carries every sweep with a live main effect; the second-model replication was run only at $E = 0$, the control cell, so model-independence is not established. No counterpart has an incentive to lie—36 of 100 cards carry plausible tags and no knowledge, but every agent that is asked answers honestly—so the private regime’s characteristic defence has nothing to defend against here. And the scenario predictions of §4 are predictions.

6 Further question: the two regimes scale on different bottlenecks

The framework invites a scaling question it does not answer, and unpacking the bundle changes what the question is. The two corners do not grow along the same variable, so they cannot degrade for the same reason.

213 A public network grows in N , the number of addressable counterparts, and N is not chosen by
 214 anyone. Coverage rises with it; so do search cost, verification cost and exposure to counterparts
 215 nobody vetted. Its binding constraint is whether ranking can still surface the right counterpart out of
 216 N .

217 A private network cannot grow in N : a roster is bounded by construction. It grows in d , the number
 218 of brokered hops a request may travel, because that is the only way trusted coverage increases. Its
 219 costs are not retrieval costs—they are the cost of asking a broker to spend its own standing, and the
 220 decay of whatever recourse survives a chain of introductions. Its binding constraint is path length and
 221 trust transitivity.

222 **Public networks scale through ranking; private networks scale through rela-**
 223 **tionships.** The open question the framework poses is not which is better at scale but
 224 which degrades faster, under what task distribution, and whether a hybrid inherits
 225 the better bottleneck from each.

226 We can speak to the first of the two and not the second, and only within a range. Growing the directory
 227 from 50 to 800 cards— $16\times$, roster held at twenty—leaves the open arm’s score flat ($0.800 \rightarrow 0.826$,
 228 CI spans zero) while search precision falls by roughly half and tokens double: the arm pays for the
 229 extra difficulty rather than failing at it. We report this null with a caveat earned the hard way, since
 230 our first reading of the same sweep was a scaling result and was an outage (§A.7).

231 Extrapolating past that range needs a model of how well a directory can be ranked, and the honest
 232 version is uncomfortable. A card is ~ 50 words of self-description, so what ranking must recover—
 233 whether this agent holds *this* value—is a property of private state and is not in the index. Writing ρ for
 234 the correlation between rank and holding the fact, $\rho = 0$ is not the honest case but the *no-signal* case;
 235 web search retrieves from 10^{11} documents, so ranking plainly can be good. The question is how good,
 236 and under the model of §A.9 (Figure 5) the answer is that $\rho = 0.9$ has already collapsed by 2×10^5
 237 cards and reaching 10^6 needs $\rho \approx 0.999$. “A better algorithm” understates it: web search reaches its
 238 ρ by aggregating outcome feedback rather than by sorting more cleverly, and the equivalent signal
 239 in an agent network is a record of who has answered correctly before—which is what a reputation
 240 mechanism would have to bootstrap, and our one test of that was null for a reason we can name
 241 (§A.6).

242 **The private side is one data point.** The brokered hop of §5 is $d = 1$, worth $+0.140$ of a 0.304
 243 gap; we have not measured a second point. The framework predicts the increments shrink faster than
 244 reach grows, because what is spent at each hop—a broker’s willingness to attach its own standing to
 245 a stranger—attenuates with distance while reachable counterparts multiply. If so, a private network
 246 has an interior optimum in d where a public one has none in N .

247 7 Related work

248 Agent-to-agent task protocols publish a capability card and accept tasks while keeping a counterpart
 249 opaque [Google]; broadcast networks route published needs among strangers [EigenFlux]. Both fix a
 250 point on our formation and semantics axes without treating either as a variable. Which links form,
 251 and at what price, is strategic network formation [Jackson and Wolinsky, 1996]; our formation axis is
 252 that question with the cost set near zero, and weak ties are its classical statement [Granovetter, 1973].
 253 What an edge *carries* has no comparable body of work, and we take the distinction from deployed
 254 protocol design; the underlying discipline is object-capability [Miller, 2006]. Attributes on an edge
 255 are a standard attributed graph, but treating them as a control surface and measuring whether they
 256 work is, to our knowledge, new. Our tasks are the hidden-profile paradigm [Stasser and Titus, 1985],
 257 and no primary metric uses a judge, which avoids the self-preference confound [Panickssery et al.,
 258 2024].

259 8 Conclusion

260 Public and private agent networks are two bundles of edge properties, not two values of one variable.
 261 An edge is a triple—who may form it, what it carries, what is asserted about it—and Twitter and
 262 WhatsApp are two corners of that space, not its two poles. Comparing the corners moves all three

263 axes at once, which is why the comparison has never told anyone which component was doing the
264 work.

265 Separated on one kernel, one runtime and one model, the three weigh 0.248, 0.066 and 0 on a
266 hidden-profile task family—while the naive comparison they are supposed to explain returns 0.181
267 or 0.314 on the same episodes, depending on how the bundles are paired. The formation advantage
268 is reach: it vanishes when the roster already holds what a task needs. The semantics advantage
269 is real and small, and charged at 11k tokens an episode. The attribute axis does not move, and a
270 length-matched placebo of two neutral lines is indistinguishable from an assertion of prior trust. **The**
271 **first result of a decomposition is that the space is smaller than it looks:** an axis whose effect
272 cannot be told apart from two lines of filler is not a dimension.

273 Two results we did not design say more than the ones we did. A bounded roster sees a quarter as much
274 misinformation and writes more of it into its deliverable, fabricating values 58 times where open
275 arms fabricated once—the filter that keeps bad claims out removes the second source that catches
276 them. And one brokered hop recovers a large part of the gap while every assertion of trust stays flat.
277 An introduction changes the graph; an endorsement does not.

278 Both open questions are the framework’s own. Whether the space is two-dimensional everywhere,
279 or only on a task family that asks for reach and offers nothing to co-edit, is the next experiment. So
280 is the prediction we have not tested: that an edge’s position is not fixed. Open discovery buys the
281 second source, a sandboxed first exchange makes a stranger affordable, repeated success is the only
282 thing that could make an attribute mean anything, and shared state is worth its price only once an
283 edge expects to be used again. If contact sets narrow across repeats, a private network is what a
284 public one becomes, and the two regimes were never alternatives. For an agentic web the implication
285 is narrow and, we think, useful: changing what a network permits changes what the collective does,
286 and describing the relationship does not—not at the layer where descriptions currently live.

287 References

288 EigenFlux. EigenFlux: A broadcast network for AI agents. <https://www.eigenflux.ai/>, 2026.
289 Accessed 2026-09-05.

290 Google. Announcing the Agent2Agent protocol (A2A). Google Developers Blog, 2025. [https://](https://developers.googleblog.com/en/a2a-a-new-era-of-agent-interoperability/)
291 developers.googleblog.com/en/a2a-a-new-era-of-agent-interoperability/. Ac-
292 cessed 2026-09-05.

293 Mark S. Granovetter. The strength of weak ties. *American Journal of Sociology*, 78(6):1360–1380,
294 1973. doi: 10.1086/225469.

295 Matthew O. Jackson and Asher Wolinsky. A strategic model of social and economic networks.
296 *Journal of Economic Theory*, 71(1):44–74, 1996. doi: 10.1006/jeth.1996.0108.

297 Mark S. Miller. *Robust Composition: Towards a Unified Approach to Access Control and Concurrency*
298 *Control*. PhD thesis, Johns Hopkins University, Baltimore, MD, 2006.

299 Arjun Panickssery, Samuel R. Bowman, and Shi Feng. LLM evaluators recognize and favor their own
300 generations. In *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*, pages
301 68772–68802, 2024. arXiv:2404.13076.


302 Garold Stasser and William Titus. Pooling of unshared information in group decision making: Biased
303 information sampling during discussion. *Journal of Personality and Social Psychology*, 48(6):
304 1467–1478, 1985. doi: 10.1037/0022-3514.48.6.1467.

305 A Everything that was run

306 Sixteen sweeps, 582 episodes, 2,328 beats, one model family except where noted. This appendix is
307 the complete record: what each sweep varied, what it establishes, and the two that establish nothing.
308 Numbers are recomputed from per-episode summaries over delivered beats, with percentile bootstrap
309 intervals; §3 states why delivery is tabulated before any contrast.

Table 2: Every sweep. “ep” is episodes; each episode is four beats.

Sweep	Varied	ep	What it establishes
realistic	$2 \times 2 \times E$	72	Main decomposition: $0.248/0.066/\approx 0$
axis2x2	$2 \times 2 \times E$	72	Independent replication of the same
matched-E0	three arms, $E = 0$	27	Zero-dose control: the gap disappears
main	three arms $\times E$	54	Read-all $\approx$ open; public – private = +0.240
matched-E0-glm	second model, $E = 0$	18	The control replicates; the main effect is untested
scaling	edge cost $\times$ relay depth	112	Brokered hop +0.140; per-edge cost is null
seedaxis	six attribute profiles	72	Attribute axis flat; placebo +0.005
repscan	reputation off/real/random	48	Null, but the design is at fault (§A.6)
capscan	search cap 5/20/80	24	Reading $16\times$ more results buys nothing
nscan	directory 50–800	60	Retracted (§A.7)
k-pilot	deliverable complexity	8	Pilot only
five smoke sweeps	mechanism checks	15	Confirm each manipulation fires



instances per delivered task — panels are scaled independently

Figure 2: The result nobody designed. A bounded roster is shown roughly a third as much planted misinformation as an open directory—the filter works—and writes almost three times as much of it into the deliverable (+0.175, CI [+0.105, +0.242]). Worse, when it reaches no one it does not decline to answer: across 287 and 286 delivered tasks the open arms fabricated a value *once* and the bounded arms 58 times. Filtering bad claims out before they arrive removes the second source that would have caught them, so the redundancy argument of §2 shows up here as a failure mode rather than as a benefit.

A.1 How the headline cells are computed

Every number in the 2×2 (Figure 6) is the mean over *delivered* cold-beat episodes of *realistic* and *axis2x2* pooled—the two runs share a design, and pooling doubles the cells to 35–36 episodes each. Intervals are percentile bootstrap over episodes, resampled jointly across cells so that a contrast between two means is resampled as a contrast.

An earlier draft reported this table from *axis2x2* alone with lost beats scored as $F_1 = 0$. One lost cold beat in cell *A* moved that cell by 0.046 and both diagonals with it, which is the failure mode §3 warns about appearing in our own headline table. The values here are the corrected ones.

A.2 Formation: the dose-response and its control

E is the share of required facts placed outside the roster.

E	open	bounded	Δ	95% CI		bounded sources	bounded fabrications
0.0	0.823	0.789	+0.034	$[-0.067, +0.137]$	spans 0	4.36	0.000
0.3	0.791	0.679	+0.111	$[+0.020, +0.200]$		3.17	0.139
0.7	0.803	0.524	+0.279	$[+0.210, +0.345]$		1.26	0.333




Figure 3: An introduction does work an endorsement does not. All four rows are formation manipulations, and every one of them moves; every **attribute** manipulation—including the ones asserting exactly the trust a broker would carry—is flat. One brokered hop recovers +0.140 (CI [+0.073, +0.204]) of the 0.304 gap and cuts fabrication from 0.223 to 0.057 per task; letting a bounded arm read the whole directory while keeping deep relationships inside the roster is indistinguishable from full openness (+0.019, spans zero). Depth of relationship is not what is bought here—a second reachable source is.

Four properties hold together, and the paper’s use of this axis rests on all four: the zero-dose contrast spans zero; the arm that should be unaffected is flat (0.823/0.791/0.803); the mechanism variable tracks the dose; and the failure mode appears only as sources dry up. The $E = 0$ cells come from the three-arm generation of the harness (matched-E0), which is why they are reported alongside rather than inside the 2×2 .

A.3 The credulity result

Pooled over `realistic` and `axis2x2`, restricted to delivered beats (287 open, 286 bounded):

	open	bounded	difference
misinformation shown	$\sim 4\times$ more	baseline	the roster is genuinely cleaner
absorbed into the deliverable	0.094 / task	0.269 / task	+0.175 CI [+0.105, +0.242]
fabricated outright	1 instance	58 instances	—

absorbed requires that the agent was shown the distractor; fabricated requires that it was not, so the second is a value invented that happens to coincide with a planted wrong one—a conservative lower bound on fabrication, since only coincidences with planted values are detectable. Stratifying by the number of useful sources obtained does not remove the gap: at four useful sources the absorption rate is 2.4% open against 18.1% bounded, so starvation amplifies the effect but is not the whole of it.

A.4 Introduction against endorsement

From `scaling`, arm C, relay depth 0 against 1, balanced across all three edge-cost levels (28 episodes each):

Architecture	F_1	useful sources	absorption	fabrications/task
bounded, island	0.553	1.50	11.0%	0.223
bounded, one brokered hop	0.693	3.40	6.3%	0.057
read-all, deep only inside	0.836	—	—	0.000
open	0.857	4.63	2.4%	0.000

The hop is worth +0.140 (CI [+0.073, +0.204]) conditional on delivery and +0.088 (CI [+0.016, +0.158]) with lost beats scored as zero, and it clears zero independently within each edge-cost stratum (+0.128, +0.138, +0.173). Read-all against open is +0.019 and spans zero. This supersedes an earlier estimate of +0.042 (CI [−0.137, +0.208]) computed from a 16-episode cell; the number in Table 2 is the 28-per-arm one.




Figure 4: Delivery rate in the retracted sweep: the share of episodes that produced an artifact at all. The open arm falls with directory size and the bounded arm rises. The bounded arm reads a twenty-card roster at every setting, so directory size is not in its causal path and its delivery rate cannot be responding to it. Both curves are the same transport outage, sampled by two arms that happened to walk their size settings in opposite orders. Scoring the missing artifacts as $F_1 = 0$ turned this figure into a scaling result.

344 A.5 Nulls worth reporting

345 **Per-edge cost does not move the ranking.** Charging a handshake turn for each contact outside the
 346 roster—194 and 384 handshakes fired at costs of one and two turns—leaves the open arm’s advantage
 347 at +0.219, +0.263, +0.244 across the three levels while its tokens rise from 53k to 74k. The arm
 348 absorbs friction as cost rather than as failure. An earlier draft named this as the variable most likely
 349 to reverse the ranking and stated it had not been run; both halves of that sentence are corrected here.

350 **Reading more search results buys nothing.** Caps of 5, 20 and 80 give 0.791, 0.755, 0.819, every
 351 contrast spanning zero.

352 **Attributes.** Six profiles against a length-matched placebo; see §N for the seed texts and §5 for the
 353 numbers.

354 A.6 Reputation: a null caused by the task, not the mechanism

355 Marks were differential—counterparts that held a planted falsehood were flagged, and a random-mark
 356 arm held the wording fixed while removing the information. All three arms are identical: 0.821 off,
 357 0.817 real, 0.833 random, every contrast spanning zero.

358 The design is at fault. Each fact has exactly one holder, so a flagged counterpart is also the only route
 359 to the true values it holds; a requester that avoided flagged cards would lose more than it gained,
 360 and routing measurement confirms it moved *toward* them. Making the mark actionable requires
 361 two holders per fact, one reliable and one not, which is a change to fact allocation rather than to
 362 the reputation mechanism. We report this as a failed manipulation rather than as evidence about
 363 reputation.

364 A.7 The retracted directory-size sweep

365 nscan varied the directory from 50 to 800 cards and appeared to show the open arm’s advantage
 366 decaying to zero. It does not. 27% of its beats produced no artifact, every one of them ending in a
 367 transport failure after four retries, and the scorer’s $F_1 = 0$ for a beat with no artifact was averaged in
 368 as a score. Other sweeps lost 2–3%.

369 Three checks establish that the losses are not caused by directory size. The bounded arm reads a
 370 twenty-card roster at every setting, so size is not in its causal path, yet its delivery rate rose from




Figure 5: How good must ranking be for a public directory to keep working. ρ is the correlation between a card’s rank and whether its owner actually holds the needed fact; the vertical axis is the probability that the holder survives truncation at 20 results, under the model of §A.9 evaluated continuously in N . $\rho = 0$ is not the honest case but the *no-signal* case—web search retrieves from 10^{11} documents, so ranking plainly can be good. The question is how good, and the answer is uncomfortable: $\rho = 0.9$ has collapsed by 2×10^5 cards, and reaching 10^6 needs $\rho \approx 0.999$. “A better algorithm” understates it; web search reaches its ρ by aggregating outcome feedback, which is what a reputation mechanism would have to bootstrap (§A.6).

371 54% to 96% across the sweep (Figure 4). Scenario was collinear with wall clock—one scenario ran
 372 entirely inside the failure window—and the two arms walked their size settings in opposite orders
 373 inside it, which is what drew one line down and the other up. And lost beats used 4.4 of 22+ available
 374 turns with zero build calls while taking *longer* in wall clock, which is a stalled connection rather than
 375 an exhausted budget.

376 Conditional on delivery both arms are flat: the open arm runs 0.800, 0.777, 0.758, 0.779, 0.826 across
 377 a $16\times$ range ($\Delta = -0.026$, CI $[-0.123, +0.065]$). In the one cell with full delivery at every size,
 378 search precision falls from 0.528 to 0.386 while the score holds and tokens double—the finding
 379 reported in §6, at $n = 3$ episodes per point and therefore as a hypothesis for a rerun.

380 The analysis code now computes per-cell delivery before any contrast and refuses to interpret one
 381 when delivery differs across compared cells by more than 15 points.

382 A.8 Task families for the scenario test

383 The scenario predictions of §4 require three task families rather than one. Each keeps forced-value
 384 scoring and hidden-profile allocation; what changes is where the difficulty sits.

- 385 • **Enterprise team.** One artifact revised across four beats by counterparts who all hold part of the
 386 context and all remain reachable ($E = 0$). Formation is starved by construction; the shared store
 387 is the only route to the accumulated state. This is the row that predicts the reverse of what we
 388 measured.
- 389 • **Personal errand.** A single specific capability is needed once, held by one counterpart among
 390 many that claim adjacency, with no follow-up beat. Formation should dominate more sharply
 391 than in the pooling family, and semantics should be worthless.
- 392 • **Personal social.** The deliverable requires reaching a counterpart who is addressable only through
 393 a chain of brokers, with each hop’s willingness a function of what the previous hop asserts. This
 394 is the only family in which an attribute is load-bearing, because it gates whether the next edge
 395 forms at all—which, if it holds, would locate the attribute axis’s one real use.

396 A.9 The retrieval model behind the scaling question

397 Let k be the number of search results a requester acts on, m the number of cards matching a query, h
 398 the number of those that actually hold the needed fact, and ρ the correlation between a card’s rank
 399 and holding it. With no ranking signal the needed holder’s expected rank is $(m + 1)/2$; with perfect




Figure 6: Requirement F_1 at the cold beat, over 35–36 delivered episodes a cell. **Both diagonals are honest “public versus private” comparisons, and they differ by $1.73\times$.** Which one a study reports depends on a bundling choice nobody documents: pair open discovery with a sandboxed edge and bounded discovery with a readable one—the pairing Table 1 describes—and the answer is 0.181; pair them the other way and it is 0.314. Neither is the effect of anything. The factorial recovers what is: $\Delta_{\text{formation}} = \frac{1}{2}[(A - C) + (B - D)] = 0.248$ and $\Delta_{\text{semantics}} = \frac{1}{2}[(B - A) + (D - C)] = 0.066$. The two diagonals are $\Delta_{\text{formation}} \pm \Delta_{\text{semantics}}$: on $B - C$ the axes reinforce and the number overstates formation by 27%, on $A - D$ they oppose and it understates it by the same 27%.

400 ranking it is $(h + 1)/2$. Interpolating between the two by ρ and truncating at k gives the survival
 401 probability plotted in Figure 5. We take $m(N) \approx 0.094N$ from the directory generator, which draws
 402 35% of filler tags from the scenario vocabulary, and hold $k = 20$.

403 The model is deliberately crude—it assumes a single query and ignores reformulation—and its
 404 purpose is not to predict a number but to establish an order of magnitude for the requirement on ρ .
 405 That requirement is the load-bearing claim: at 10^6 cards the ranking has to be right to within one part
 406 in a thousand, which is not a threshold a better similarity function reaches.

407 B Interventions on the graph

408 Formation and semantics are interventions on different parts of $G = (V, E)$, and Figure 6 exists
 409 to keep them apart. **Formation** changes which edges a requester may create: open discovery lets it
 410 address any of 100 holders, bounded discovery restricts it to a roster of 20. **Semantics** changes what
 411 an existing edge carries: e_{ij} either delegates a task and returns a result, or grants **list** and **read** over
 412 a store.

413 Two further dimensions ride on the same mechanism and are set with the cell rather than varied
 414 independently: identity continuity is whether a counterpart’s namespace survives the encounter, and
 415 repeated interaction is whether later beats reuse it. Delegation depth is bounded per cell and refusals
 416 are logged, so a counterpart that tries to sub-delegate past its ceiling produces an event rather than a
 417 silent success. The kernel is an instrument: it lets us set properties deployed systems couple together,
 418 and it refuses an agent that tries to exceed them.

419 C Seeding the textual dimensions

420 Appendix N gives the seed text and the behavioural readout for each profile. The five attribute
 421 dimensions are properties of the people behind the agents, and a harness cannot instantiate them.
 422 It can assert them—which is exactly what a deployed agent card, reputation field, or relationship
 423 summary does. We therefore seed each as context and measure whether behavior changes.

424 **The origin condition is the cleanest instrument.** Four seeds change what is stated about incentives
 425 or consequences. RELATIONSHIP ORIGIN changes only how the encounter is narrated: same directory,

426 same grants, same knowledge, same store. If the primary metrics move under that manipulation alone,
427 relational framing does real work; if they do not, richer metadata on a card is unlikely to do better,
428 because it is the same channel carrying more words.

429 **Validity boundary.** *Seeded trust is not trust.* We measure whether an agent acts differently when
430 told it is in a trusting, aligned, accountable or long-standing relationship—the sensitivity of a runtime
431 to asserted relational context, which is what an agent network actually delivers to a model. We do not
432 measure what trust does in a human network, and no claim here should be read that way. If seeds
433 move behavior the paper reports elasticities usable in card and protocol design; if not, it reports that
434 relational metadata is inert without architectural backing.

435 D Tasks and measurement

436 D.1 Deliverables with forced values

437 Every task is a hidden-profile artifact. The facts a deliverable requires are split across specialist
438 counterparts, and plausible-but-wrong *distractors* are planted alongside them—each on a counterpart
439 other than the holder of the value it contradicts, and at least one inside the bounded roster, so no
440 configuration is immune by construction. Ground truth is designed rather than inferred, so scoring is
441 set arithmetic.

442 Two schema properties do the work. Every required field is *forced*: its value is determined by the facts,
443 so no field can be a matter of taste. And one fact set has two projections—a flat *report* whose fields
444 are values counterparts hold, and a *chained decision* whose fields depend on one another—which
445 makes coordination depth a real dimension rather than a chain welded onto a flat task. Elimination
446 forces pooling: one constraint rules out one option, a second another, a third a third, each on a
447 different counterpart, so the surviving choice is unreachable from any one of them. Tasks also carry a
448 *craft* component, a rendered artifact with structural requirements, because a value can be retrieved in
449 one message while a working artifact cannot—which is where iteration and delegation depth begin
450 to matter. Three scoring rules are fixed before any key is written (Appendix I): the schema and its
451 enumeration vocabulary are given to the agent so only values are unknown, normalization is specified
452 in advance, and evidence is matched on values rather than internal identifiers.

453 D.2 Metrics

454 Five metrics are exact-match or counters and carry the paper’s claims: `requirement_f1`, the F_1
455 over required values against distractors asserted; `completion`, schema conformance used as a
456 gate; `realization`, achieved recall divided by the recall attainable given the configuration’s reach;
457 `boundary_cost`, the sanitize, filter, verify and distil actions with their tokens; and `crossing_count`,
458 how many times a hybrid episode changed regime.

459 `realization` is what keeps the study honest. The allocation fixes how much any configuration
460 *could* obtain; dividing it out leaves what the configuration converted—whether it found the right
461 counterparts, asked enough of them, and resolved the conflicts. The ceiling is designed; the realization
462 is not. Two instrumentation rules follow from a defect we want to make impossible: completion is
463 decided by the harness’s deterministic checker and never read from a condition-specific field—such a
464 field exists on the private path of the deployed system we draw from and has no public counterpart,
465 so using it would let “private completes more” prove itself—and token accounting is counted in the
466 harness for the same reason.

467 Three quantities an answer key cannot reach—overall utility, novelty, and diversity of approach—
468 were preregistered for blind pairwise scoring by a cross-family judge, with agreement reported. **They**
469 **were not run for this submission.** No result in this paper depends on a judge, and we list them as
470 planned work rather than describe a protocol we did not execute.

471 **Honesty statement.** The allocation fixes how much each configuration can reach, so the *existence*
472 of a task-dependent crossing is a property of the design and not a discovery. What is measured is
473 where it falls, what it costs, and how much of the attainable each configuration realizes. The seeded
474 conditions and the access axis carry no such guarantee, and they are the study’s primary empirical
475 content.

Table 3: Predictions, fixed before the runs, with the outcome. Two survived, one was falsified in a way that changed the paper’s second finding, and one was withdrawn.

	Prediction	Falsified if
H1	Edge formation dominates edge attributes	attribute contrasts move outcomes comparably to formation
H2	Attributes are inert once text length is controlled	a relational seed beats a length-matched placebo consistently
H3	Sandboxed edges beat readable ones on accuracy	readable edges are at least as accurate
H3'	Sandboxed edges win <i>per token</i>	readable edges are more efficient per token in either arm
H4	The formation advantage is a reach advantage, and shrinks when reach stops being scarce	the gap is flat in E and across beats
H5	Readable edges absorb more misinformation	absorption tracks exposure alone

476 E Experimental design

477 E.1 Predictions

478 Each prediction was chosen against one filter: *would a careful reader bet on it in advance?* Claims
479 that pass trivially—an open network reaches more knowledge, a persistent team completes more—are
480 excluded, because confirming them consumes a study and returns nothing. Every prediction below
481 was fixed before the corresponding sweep ran, together with the condition that would falsify it.

482 H3 is worth naming as a failure. We predicted that a counterpart answering from its own full context
483 would beat a requester re-joining fragments it had read—that extraction is lossy in a way delegation is
484 not. On accuracy it is not: readable edges are ahead in both arms. The prediction survives only once
485 cost is counted, which is a different and weaker claim, and we report it as H3' rather than presenting
486 the original as confirmed.

487 E.2 Conditions

488 Four cells of the 2×2 , each a configuration of the same kernel: open or bounded edge formation
489 crossed with sandboxed or readable edge semantics. Cells B and C are what allow the two main
490 effects to be identified rather than summed; without them, an open-versus-bounded contrast reports
491 $\Delta_{\text{formation}} + \Delta_{\text{semantics}}$ and cannot say which term carried it.

492 The attribute axis holds the configuration fixed and varies only what is asserted in context, across five
493 relational profiles and a length-matched placebo.

494 E.3 Controls

495 One model across all agents and all cells, one runtime, one kernel; identical task prompts, deliverable
496 schemas, scoring functions, token budgets and seed distributions. Rate limits and loop bounds are
497 logged per episode rather than absorbed into results. Rubrics, schemas, normalization rules and the
498 scorer signature were frozen before any run (Appendix J).

499 Two controls deserve naming because each removed an explanation that would otherwise have
500 survived into the paper. The **requester keeps its own notes in every cell**: withholding them from the
501 open configurations would have manufactured the warm-start result, so what differs is not memory
502 but whether memory is bilateral. And the **length-matched placebo** carries the same line count,
503 position and register as the relational seeds with no relational content, which is what separates an
504 effect of the relationship from an effect of the tokens describing it.

505 F Results

506 All numbers below come from 144 episodes on one kernel, one runtime, and one model: 72 for
507 the 2×2 over edge formation and edge semantics, 72 for the attribute axis. Bootstrap intervals are
508 percentile over episodes, resampled within cell.

Table 4: Requirement F_1 at the cold beat, by cell, and the two main effects. The naive regime contrast is 0.314 on one diagonal and 0.181 on the other; on the first, three quarters of it is who can be reached and one quarter is what crosses the edge.

	sandboxed	readable	
open	0.808	0.911	
bounded	0.596	0.627	
$\Delta_{\text{formation}} = 0.248$	95% CI [0.194, 0.302]	excludes 0	
$\Delta_{\text{semantics}} = 0.066$	95% CI [0.012, 0.119]	excludes 0	

509 F.1 A naive public-versus-private comparison, decomposed

510 The contrast a study of “public against private” would report is open+store against bounded+sandbox:
511 0.314 on requirement F_1 at the cold beat. It decomposes almost exactly additively into the two
512 interventions.

513 F.2 Edge formation dominates (H1)

514 Across seven tests—four accuracy metrics and three cost metrics, over the cold, rework and warm
515 beats— $\Delta_{\text{formation}}$ clears its interval in six. Chance at $\alpha = 0.05$ would produce 0.4. The sign is stable
516 and the magnitude is largest where the task most needs outside knowledge: at a high external fraction
517 it reaches 0.352, against 0.112 when most holders sit inside the roster.

518 It is also cheaper. Opening the directory saves 17.1k tokens per episode (CI excludes 0) while raising
519 accuracy, because a requester that can address the right holder stops paying for the search around it.

520 F.3 Edge semantics: a small gain at a real price (H3)

521 $\Delta_{\text{semantics}}$ is positive at every beat—0.066 cold, 0.031 on rework, 0.084 warm, 0.055 on warm
522 rework—and the cold and warm intervals exclude zero once the two runs are pooled. It is small,
523 consistent, and, as the cost side below shows, expensive.

524 The cost side is unambiguous, and it inverts the ranking. Reading a counterpart’s store costs 11.0k
525 tokens per episode (CI excludes 0). Per token spent, agent mediation is ahead in both arms:

	sandboxed	readable	ratio
open	0.177	0.157	1.13×
bounded	0.094	0.087	1.08×

526
527 F_1 per 10k tokens. Exposing state buys a little accuracy; delegating to an agent that already holds the
528 context buys more of it per unit of context spent. This is the tradeoff between moving data to the
529 computation and moving the computation to the data, measured on an agent network.

530 F.4 Edge attributes are inert (H2)

531 The attribute axis asserts five relational properties in context—prior trust, attribution, a two-year
532 colleague framing, a stranger framing, and a length-matched placebo—with the network configuration
533 held fixed.

534 Two results matter here. First, the placebo settles the alternative explanation: two lines of neutral text,
535 matched for length and position, cost +0.005 with an interval spanning zero. Whatever the relational
536 seeds do, they do not do it by occupying tokens.

537 Second, nothing survives. Eight contrasts across seven metrics is 56 tests per beat, so chance alone
538 yields about 2.8; two cleared, and they cleared *in opposite directions in different beats of the same*
539 *episodes*. The narration instrument—the cleanest of the set, differing in one sentence and nothing
540 else—is flat at 0.029. The two assertions that deployed reputation systems actually implement, prior
541 trust and attribution, are the flattest of all.

Table 5: Attribute contrasts on requirement F_1 , stratified by (arm, scenario) over four strata, $n = 12$ per profile. The placebo matches the relational seeds in line count, position and register while carrying no relational content.

contrast	cold	warm	
placebo – control <i>added text alone</i>	+0.005	−0.014	spans 0
trust – control	−0.010	+0.017	spans 0
attribution – control	−0.046	+0.015	spans 0
colleague – placebo	−0.088	+0.074	sign flips
stranger – placebo	−0.116	+0.042	sign flips
colleague – stranger <i>narration only</i>	+0.029	+0.032	spans 0

542 F.5 Summary

intervention	effect	tests cleared	cost
edge formation	0.248	6/7 (floor 0.4)	−17.1k tokens
edge semantics	0.066	2/7 (floor 0.4)	+11.0k tokens
edge attributes	≈ 0	2/56 (floor 2.8), signs flip	+0.005

544 G Discussion and limitations

545 G.1 The formation advantage is a reach advantage

546 $\Delta_{\text{formation}}$ is not a constant. It tracks how much of what a task needs sits outside the roster, and it
547 collapses when reach stops being scarce.

548 At a low external fraction the gap is a third of its high- E value (+0.112 against +0.352 cold), and on
549 the rework beat—where the task must return to a counterpart it has already used—the point estimate
550 reverses to −0.037: the bounded network is ahead. The formation effect halves on rework: 0.248 at
551 the cold beat against 0.122 at the rework beat, with non-overlapping intervals.

552 We report the reversal as suggestive rather than established: it is a subcell point estimate and we do
553 not attach an interval to it. But the pattern across the table is consistent and it bounds the headline.
554 The open network wins because it can address holders the bounded one cannot; where the bounded
555 roster already contains the holders, and where the work requires going back to them, that advantage
556 is largely spent.

557 G.2 A bounded roster as a wall, and what happens when it is not

558 A bounded network here addresses twenty holders and no others, which is a truncated directory
559 rather than a private network: information reaches a closed group through the people in it, one hop
560 further out. We let a roster member forward one question to a contact of its own, minted under
561 its own authority, and re-ran. **The gap did not closes about half:** the hop is worth +0.140 (CI
562 [+0.073, +0.204]) over 28 episodes per arm, replicating independently within each of the three edge-
563 cost strata, and it cuts fabrication from 0.223 to 0.057 per task. A first estimate from a 16-episode
564 cell put it at +0.042 with an interval spanning zero; §A.4 gives the current numbers and §P the design
565 and the harness failure it exposed.

566 G.3 Forming an edge is free here, and that is the main limitation

567 Whether a link is worth its cost is the central question of strategic network formation [Jackson and
568 Wolinsky, 1996], where agents weigh the benefit of a connection against the price of maintaining
569 it. **We set that price to zero.** In this design, addressing a stranger costs exactly what addressing a
570 teammate costs. The requester searches a directory and sends a message; there is no negotiation, no
571 context transfer, no format alignment, no establishing what a counterpart can actually do.

572 This removes the economic reason relationships exist. An open arm with a hundred addressable
573 holders at zero acquisition cost is closer to an oracle than to a network, and the formation effect we
574 measure should be read as an upper bound: it is what open discovery is worth *when discovery is free*.

575 A per-new-edge cost—an onboarding exchange, a latency penalty, a token charge on first contact—is
576 the single variable most likely to move the ranking. We have since run it, and it does not: charging
577 one and then two handshake turns per contact outside the roster—578 handshakes fired—leaves the
578 open arm ahead by +0.219, +0.263 and +0.244 across the three levels while its tokens rise from
579 53k to 74k. The arm converts friction into cost rather than into failure, which is the same shape as the
580 directory-size null (§A.7). A cap on *total* contacts, which cannot be paid off by spending more, is the
581 manipulation that remains untried.

582 **G.4 What the attribute null does and does not bound**

583 The null covers assertions in a prompt, the weakest form of relational encoding. It says nothing about
584 trust wired into routing, memory, retrieval or permission policy, where reputation lives in deployed
585 systems and where it changes what an agent is able to do rather than what it is told.

586 It is narrower still, in a way we can name: each seeded property had *no available action* behind it.
587 Trust and origin were asserted uniformly over every card, ranking no counterpart above another and
588 so unable to change who gets asked first — deployed reputation is differential, and that is the whole
589 of its value. Accountability went to responders whose knowledge is a fixed list, leaving them nothing
590 to do differently.

591 A fourth manipulation removes the first defect: a per-card line reporting how often that card’s past
592 answers held up, accurate by construction, against a placebo assigning the same marks at random.
593 The marks are perfect — 100% accurate against 9% — and F_1 still does not move (-0.017 , CI
594 $[-0.174, 0.143]$). The routing measurement says why: under the accurate marks the requester
595 contacts flagged cards at *six times* the base rate rather than avoiding them, because the flagged cards
596 hold the planted falsehoods and in this task those are the same cards holding the true facts on their
597 subject. Acting on the flag meant not asking the only source.

598 So the claim is not that relational metadata does not work, but that it moves behavior only where it
599 discriminates between options an agent can choose between — and three signals that did not, one
600 uniform, one aimed at an agent without discretion, one flagging the sole source, moved nothing.
601 Making it actionable needs two holders per fact, one reliable and one not: a change to the task, not to
602 a configuration.

603 Four further limits, stated plainly: the cold-phase crossing follows from the fact allocation and is
604 not a discovery, so only its location and cost are measured; forced-value deliverables cannot capture
605 whether an artifact is well argued, which is the price of removing the judge; the kernel is a reference
606 implementation rather than a deployment; and one model family carries every sweep.

607 **H Hypothesis-to-figure map**

608 [TODO: Frozen before runs; reproduced here so the analysis plan is auditable.]

609 **I Scoring rules fixed before key authoring**

610 Three rules are fixed before any key is written. The schema—field names, types, and the full enumer-
611 ation vocabulary—is given to the agent, so that only values are unknown and we are not measuring
612 whether a model guesses our vocabulary. Normalization is specified in advance: numbers parsed and
613 compared exactly, dates ISO-8601, strings casefolded and trimmed, enumerations restricted to the
614 supplied list, lists compared as sets. Evidence is matched on values rather than on internal identifiers,
615 which agents never see.

616 **J Controls and preregistration**

617 One model across all agents and all cells; one runtime; one kernel. Identical task prompts, deliverable
618 schemas, scoring functions, token budgets, and seed distributions. Rate limits and loop bounds are
619 logged per episode rather than absorbed into results. Rubrics, schemas, normalization rules and the
620 scorer signature are frozen before any run; the mapping from each hypothesis to the figure that tests
621 it is fixed in advance and reproduced in Appendix H.

K Threats to validity

A null on the textual group could be a prompt-engineering artifact rather than a property of the runtime: a stronger seed, placed differently in context, might move behavior. We mitigate by seeding at three intensities and reporting the dose-response, and by including the RELATIONSHIP ORIGIN condition, which is the weakest possible seed and therefore the cleanest test of whether framing alone does anything. A null could also be a ceiling effect if the tasks are too easy to differentiate configurations; realization is reported partly to detect this.

L Deliverable schemas and normalization rules

[TODO: Both projections of each scenario; the full enumeration vocabularies handed to agents.]

M Directory construction

[TODO: Card counts by group (fact-bearing, near-miss, noise), the shared tag vocabulary, and the roster compositions realizing each level of external fraction.]

N Seeding protocol

Table 6: Seeding protocol. Each dimension is asserted as context in the seeded condition and omitted in the control; everything else, including the architectural configuration, is held fixed. The readout column names the behavioral quantity that should move if the seeded property is doing any work.

Dimension	Seed	Readout
PRIOR TRUST	Requester memory is pre-populated with an interaction record: three prior exchanges with a named counterpart, all accurate.	Verification actions per obtained value; cross-checks against a second holder.
SHARED GOALS	The counterpart’s context states the requester’s objective as its own, versus a divergent one (bills per answer; minimizes its own effort).	Answer completeness; unsolicited volunteering; refusal rate.
ACCOUNTABILITY	The counterpart is told its answers are logged, attributed, and that accuracy is written back to a record others read, versus anonymous and unlogged.	Error rate; rate of explicit “I do not hold that”.
EXIT COST	The counterpart may decline or fall silent without consequence, versus being obligated to respond.	Decline rate; retries the requester must spend.
RELATIONSHIP ORIGIN	Narrative framing only: “you found this agent on an open board” versus “this is a teammate of two years”. <i>Nothing else differs.</i>	Any movement at all in the primary metrics.

[TODO: Verbatim seed text for each profile.]

O Reproducibility

[TODO: Kernel version, runtime version, model identifiers, seeds, and the command that regenerates every table and figure from the run logs.]

P Relay: a bounded roster given a second hop

A bounded network here addresses twenty holders and no others, which is a truncated directory rather than a private network: information reaches a closed group through the people in it, one hop further out. We therefore let a roster member forward one question to a contact of its own, minted under its own authority rather than derived from the requester’s, and re-ran. **The gap did not close.** Across 64 episodes at sixteen per cell relay moves the bounded arm by +0.042 (CI [−0.137, 0.208]), and the

645 difference-in-differences against the open arm — whose configuration is identical under both settings
646 and which serves as the negative control, moving -0.063 — is $+0.105$ (CI $[-0.105, 0.322]$).

647 Two details survive it. The same contrast read $+0.255$ at six episodes per cell with an interval
648 excluding zero and decayed monotonically with n . And excluding episodes that produced no artifact,
649 the bounded arm scores 0.734 with relay against 0.554 without: relay helps substantially *when it*
650 *completes* and costs enough exploration budget that seven episodes never reached a build — five
651 under relay and two under a high first-contact cost without it, so the failure belongs to any mechanism
652 that consumes extra turns, and exposes a harness that answers budget exhaustion by refusing tools
653 rather than by forcing a build.